

DRAFT Markov State Modelling for the folding of HP35 villin headpiece

Miriam Zara
(Dated: June 12, 2026)

In this work, I compare different feature sets and dimensionality reduction methods for the construction of Markov state models (MSMs) of the folding of the HP35 villin headpiece.

I. INTRODUCTION

A. Brief description of the biological problem, knowns and unknowns

Markov State Models (MSMs) have become a standard framework for extracting long-timescale kinetics from molecular dynamics simulations and for constructing coarse-grained descriptions of biomolecular conformational landscapes [7]. Building an MSM typically involves several stages, including feature selection, dimensionality reduction, clustering, transition matrix estimation, validation and coarse-graining. In recent years, MSM methodology has evolved rapidly, with advances in state construction, uncertainty quantification, and the incorporation of machine learning for identifying collective variables and building models at scale, such as VAMPnets [9] and time-lagged autoencoders [10]. Beyond methodological developments, MSMs are increasingly integrated with experimental observables and applied to larger and more complex systems. Together, these efforts reflect a field that continues to evolve in both rigor and scope [1].

B. Defining the question

The purpose of this university project is to move the first steps into the world of Markov state modeling. The goal is to review and benchmark the classical MSM construction pipeline, with particular emphasis on feature selection and dimensionality reduction stages. Establishing a clear baseline based on conventional methods is a necessary step before evaluating more sophisticated machine-learning approaches in future studies.

C. Defining the plan: translating the question on the language of the computer (brief)

As a case study, I consider the villin headpiece subdomain HP35 Lys24-Nle/Lys29-Nle mutant (see Fig. 1), a small fast-folding protein for which extensive molecular dynamics trajectories are available. I compare two commonly used feature representations, native-contact distances and backbone dihedral angles, combined with two dimensionality-reduction methods, principal component analysis (PCA) and time-lagged independent component

analysis (tICA). The resulting four MSMs are evaluated in terms of (i) their ability to reproduce the folding kinetics of the system and (ii) their ability to provide an interpretable characterization of the metastable conformational states.

II. METHODS

A. Dataset assembly




FIG. 1: **Crystal structure of HP35** (PDB: 2F4K). Folded structure consists of three alpha-helices (residues 3-10, 14-19 and 22-32) that are connected by two loops. Aromatic and histidine residues are highlighted in orange: Phe6, Phe10, Phe17, Trp23, His27, and Phe35. Interactions among Phe6, Phe10, and Phe17 are central to the formation of the hydrophobic core. The two engineered mutations Lys24→Nle and Lys29→Nle are highlighted in black; these substitutions are known to accelerate the folding kinetics of the protein. Visualization was made with PyMol [14].

All analyses presented in this work are based on a single 300 μ s molecular dynamics trajectory of the fast-folding Lys24-Nle/Lys29-Nle mutant of the 35-residue villin headpiece subdomain (HP35), simulated at 360 K. HP35 is a prototypical ultrafast-folding protein with an experimentally measured folding time of approximately 0.7 μ s at this temperature and has become a standard benchmark system for Markov state model construction and validation [4–6]. The trajectory was originally generated by D. E. Shaw Research on the Anton supercom-

puter and reported by Piana *et al.* [2]. Their simulations predicted a folding time of approximately $1.5\,\mu\text{s}$, reasonably close to the experimental estimate of $0.73\,\mu\text{s}$ for the same mutant, although still larger by a factor of about $\times 2.5$. At the time of this work, the original trajectory generated by D. E. Shaw Research was temporarily unavailable due to maintenance of the data repository. I therefore used a preprocessed version distributed by Nagel *et al.* through their GitHub repository, which has been adopted in recent benchmark studies of Markov state modeling [4, 5]. The timeseries consists of approximately 1.5×10^6 samples, with a sampling interval of $\approx 200\,\text{ps}$.

B. Methods used for MSM construction used for analysis (in case you developed a method, it goes into Results)

Markov state models were built using the Deeptime library [8]. Since MSM construction is a multi-step process SF1, it is especially important to understand clearly the choices made at each step and validate them before moving on to the next one. The principle "garbage in, garbage out" applies here, and a mistake at an early stage can propagate and affect the final results. For this reason, in this section I describe synthetically the motivation for each step and define the methods that were investigated and compared in this work.

Feature definition and filtering The feature representations employed in this work were adopted from the preprocessing pipeline of Nagel *et al.* [4], to which I refer the reader for a detailed description. Two feature sets were considered: native-contact distances and backbone dihedral angles. Native contacts were defined from folded conformations observed in the simulation and filtered using distance and frequency criteria, yielding 42 contact-distance features. Backbone conformations were represented by (ϕ) and (ψ) dihedral angles using the angular transformation proposed by Sittel *et al.* [11] to account for periodicity. To remove weakly correlated motions unlikely to contribute to the folding dynamics, the MoSAIC filtering procedure of Nagel *et al.* [4] was applied. The final dihedral feature set consisted of 62 coordinates, with (ψ_1) , (ψ_2) , (ψ_{34}) , (ϕ_2) , (ϕ_{34}) , and (ϕ_{35}) excluded.

Dimensionality reduction The purpose of dimensionality reduction in the MSM workflow is not only to compress the information contained in the molecular descriptors, but also to *encode kinetic information into the geometry* of the reduced space. This aspect is particularly important because the reduced coordinates are subsequently clustered into discrete states (so-called *microstates*) that form the state set of the Markov chain. Since commonly used clustering algorithms such as k-means rely on Euclidean distance, conformations that

interconvert rapidly should ideally be mapped to nearby points in the reduced space. Otherwise, conformations separated by substantial kinetic barriers may be assigned to the same cluster, resulting in visible memory effects and inaccurate MSMs. Constructing low-dimensional representations that preserve kinetic information remains an active area of research. In this work, however, we restrict the analysis to the two classical linear approaches, principal component analysis (PCA) and time-lagged independent component analysis (tICA), which remain standard baselines in the MSM literature. PCA identifies directions of maximal variance in the data, whereas tICA identifies directions of maximal autocorrelation and is therefore explicitly designed to emphasize slow dynamical processes. A detailed mathematical description of both methods is provided in the Supplementary Information [A S1].

Both PCA and tICA require selecting the number of retained components, while tICA additionally requires specifying a lag time. To guide this choice, the reduced coordinates were evaluated through i) their free-energy profiles and ii) projected time series. The free energy associated with a coordinate z was estimated as

$$F(z) = -\ln p(z), \quad (1)$$

where $p(z)$ denotes the marginal probability density of the coordinate. The rationale behind this analysis is that coordinates relevant to the folding process should separate distinct metastable regions and therefore display multimodal free-energy profiles. Furthermore, if a coordinate captures transitions between metastable states, its time series should exhibit relatively long-lived plateaus separated by rapid transitions. For tICA, lag times between 100 and 1000 frames (approximately 20 – 200 ns) were considered. Since HP35 is known to fold a microsecond timescale, the autocorrelation - which tICA maximizes - should be inspected at times shorter than this timescale.

Clustering and selection of MSM lagtime The reduced coordinates were clustered using k-means with 100, 200, and 500 clusters. Cluster centres were initialized uniformly and the algorithm was run for a maximum of 100 iterations. Transition matrices were estimated using a maximum-likelihood MSM estimator at lag times τ ranging from 10 to 1000 frames ($\sim 2 - 200\,\text{ns}$). The counting scheme used was the "sliding window" scheme, which counts transitions at every time step, as opposed to the "discrete" scheme, which counts transitions only at multiples of τ [7, Chap. 4.1]. Detailed balance was enforced during the estimation of the transition matrix [7, Chap. 4.6]. The range of lagtimes investigated was chosen to remain substantially shorter than the characteristic folding timescale ($\sim 2\,\mu\text{s}$), ensuring that the process of interest remained temporally resolved.

Markovianity assessment and optimal lagtime choice was

performed through the implied-timescale convergence test. The i -th implied timescale estimated at lag τ , $t_i(\tau)$ is given by the relation:

$$t_i(\tau) = -\frac{\tau}{\ln \lambda_i(\tau)} \quad (2)$$

where $\lambda_i(\tau)$ is the i -th eigenvalue ($i \geq 2$) of the transition matrix estimated at lag time τ . It is known that the true timescales of the underlying continuous-state Markov process, t_i^* , are approximated from below ($t_i(\tau) \leq t_i^*$) and the relative error decreases as i) the discretization error decreases or ii) the lagtime τ increases. With this consideration, the implied timescale test consists in plotting $t_i(\tau)$ as a function of τ and looking for a region of approximate convergence, where the estimated timescales are approximately constant. The lagtime is selected as the minimum value in the region of convergence, as this choice minimizes the loss of temporal resolution while ensuring that the process of interest is approximately Markovian [7, Ch. 4.7].

C. Statistical analysis

MSM validation After selecting the MSM lag time based on the implied timescale convergence analysis, model quality was further assessed using an autocorrelation function test.

While a direct Chapman-Kolmogorov (CK) test provides a stringent validation of Markovianity, its application at the microstate level becomes impractical for models containing hundreds of states, as it requires comparing $T(k\tau)$ and $[T(\tau)]^k$ entry-wise, and summary statistics are needed which could be difficult to interpret.

Instead, following the approach suggested by [Chap. 2.4.2][7], model predictions were evaluated through the autocorrelation function of a physically meaningful observable. As an observable, the root-mean-square deviation (RMSD) from the crystal structure was chosen. For each microstate, a representative observable value was obtained by averaging the RMSD over all trajectory frames assigned to that microstate. The MSM prediction for the autocorrelation function was then computed from the spectral decomposition of the transition matrix,

$$\begin{aligned} c(n) &= \frac{\sum_{i=1}^N \lambda_i^n \langle \theta, \mathbf{r}_i \rangle_\pi \langle \theta, \mathbf{l}_i \rangle}{\sum_{i=1}^N \langle \theta, \mathbf{r}_i \rangle_\pi \langle \theta, \mathbf{l}_i \rangle} \\ &= \frac{\sum_{i=1}^N \lambda_i^n \langle \theta, \mathbf{l}_i \rangle^2}{\sum_{i=1}^N \langle \theta, \mathbf{l}_i \rangle^2} \end{aligned}$$

where λ_i and $(\mathbf{l}_i, \mathbf{r}_i)$ denote the eigenvalues and (left, right) eigenvectors of the transition matrix (in row-stochastic notation), π the stationary distribution, and (θ) is the vector of microstate-averaged RMSD values.

Equivalence between the two reported formula hold because of the detailed balance assumption. The predicted autocorrelation function was compared to the autocorrelation function computed directly from the molecular dynamics trajectory.

Spectral analysis The dynamical processes captured by each MSM were characterized through the spectral decomposition of the transition matrix (see Supplementary Information A S2 for details). The eigenvalues were converted into relaxation timescales, while the corresponding eigenvectors were used to identify the states participating in each relaxation process. States with eigenvector components of opposite sign exchange probability on the associated timescale, whereas the magnitude of each component reflects the degree of participation of that state in the process. For reversible models satisfying detailed balance, left and right eigenvectors contain equivalent kinetic information and are related through the equilibrium distribution.

Coarse-graining **TODO**

Transition Path Theory (TPT) analysis **maybe? maybe not?**

III. RESULTS

Sections describing different aspects of your analysis (3a-3d). All sections should be organised as follows, the text is continuous, broken only to a few paragraphs in each section:

- Question studied in the section and the goal of this calculation.
- Few sentences on how it was done, mostly referring to datasets and approaches in the Methods.
- Actual results and observations, each supported by either a i) table or ii) figure or iii) number in the text.
- You need to interpret the observation. For example, do not only write $X \rightarrow Y$, but also the meaning of this relationship with respect to your question.
- When presenting the results, refer to the respective Figure and Table, where it was shown.
- You can add more sentences for the interpretation. Please use precise phrases for the level of evidence, if the results 'suggest', 'indicate', 'show', 'demonstrate', 'support', 'corroborate'. Do not put too much 'hypothesis', and 'propose' these are synthetic phrases.
- End each section with a conclusion sentence for example, 'Taken together the results consistently




FIG. 2: Implied timescales convergence test. First three implied timescales as a function of the MSM lagtime. Colors denote the four combinations of feature set and dimensionality reduction method, while marker styles indicate the number of microstates used in the state-space discretization. The shaded region corresponds to $t_i \leq \tau$, where t_i is the implied timescale and τ the MSM lagtime. Timescales falling within this region are not resolved by the model, as the associated relaxation process occurs on a timescale shorter than a single Markov-chain transition step. **TODO: make fig more readable, maybe split into rows** **TODO: add uncertainties using bootstrapping**

show that X is larger than Y, indicating that X was an ancestor of Y'.

Dimensionality reduction: The tICA projections obtained using lag times between 100 and 1000 frames (20 – 200 ns) were qualitatively similar, suggesting that the method is robust within this range (Supplementary Figures SF2 and SF3]. Overall, dihedral-angle features appear to retain a richer kinetic structure than pairwise-distance features. In the tICA representation, more than seven dihedral components displayed bimodal or multimodal free-energy profiles and transition-like time series, whereas distance-based features exhibited such behaviour only for approximately the first three components. For PCA, the same visual analysis suggests generally worse performance compared to tICA, with fewer components displaying transition-like timeseries and bimodal free energy profiles. Based on this analysis, a tICA lag time of 1000 frames (200 ns) was selected for subsequent modelling. The final dimensionality-reduction parameters were: five tICs for the dihedral feature set, five tICs for the distance feature set, seven principal components for the dihedral feature set, and five principal components for the distance feature set. Next, I inspected the free energy projections of the data on the reduced coordinates and I also checked how the fraction of native contacts Q and the root mean square deviation (RMSD) to the crystal structure are distributed in the reduced space (Supplementary Figure SF5).

Clustering and selection of MSM lagtime: The implied timescale test indicates approximate convergence of the slowest dynamical processes for MSM lagtimes in the range of 20-50 ns as shown in Fig.2. The estimated timescales depend much more strongly on the choice of feature set and dimensionality reduction method than on the number of microstates used in the discretization (here, tried as 100, 200 and 500). For all four feature-reduction pipelines considered, the slowest implied timescale is in reasonable agreement ($\simeq 1 - 2 \mu\text{s}$),




FIG. 3: Results of the autocorrelation test: Autocorrelation function of the root mean square deviation (RMSD) from the crystal structure, as obtained from the molecular dynamics trajectory (blue) and predicted by the four MSMs (orange, green, red and purple). **TODO: add uncertainties using bootstrapping**




FIG. 4: Comparison between the slowest MSM relaxation mode and a structural folding descriptor. Top row: values of the first non-trivial eigenvector projected onto the first two components of the reduced coordinate space. Positive and negative components identify regions that exchange probability on the slowest resolved timescale of the Markov model. Bottom row: projection of the fraction of native contacts, Q , onto the same coordinates. The close correspondence between the eigenvector sign pattern and the distribution of Q indicates that the slowest kinetic process captured by the MSM is associated with the folding-unfolding transition.

whereas larger discrepancies are observed for the higher-order timescales. The pipeline dihedrals-tICA consistently yields the largest timescales, followed by distances-tICA, indicating more accurate approximation of the underlying continuous dynamics.

Validation Validation was performed as described in Methods comparing the autocorrelation function of the RMSD to the crystal structure computed from the trajectory and predicted by the MSMs (see Fig.3). The dihedrals-tICA, distance-tICA and distances-PCA pipeline yield similar results, while the pca-dihedrals performs substantially worse.

Spectral analysis Next, I inspected the leading left or right? check again for a row-stochastic T eigenvectors to characterize the slow relaxation modes encoded by the

Markov model. The first non-trivial eigenvector exhibits a clear separation between regions associated with low and high values of the fraction of native contacts Q , as shown by comparison with the projection of Q onto the first two tICs (Fig. 4). This indicates that the slowest dynamical process captured by the MSM corresponds to the folding-unfolding transition. The qualitative structure of the first eigenvector is remarkably consistent across all four model constructions, suggesting that the dominant kinetic process is robust with respect to the choice of molecular representation and dimensionality reduction method. In contrast, the higher-order eigenvectors display more method-dependent patterns, and their physical interpretation is less straightforward.

Coarse-graining **TODO** I build a macrostate model applying PCCA++ [Ch. 2.5.1 and 2.5.2][7].

IV. DISCUSSION

Here you try to put your results into a broader context and present your scientific view/model based on that. This is a continuous text, broken into paragraphs. In case you have two distinct 'huge' conclusions, you may generate separate sections (max 2). Chapter 4 should be organised as follows:

- You may phrase your biological question again.
- Briefly summarise the results (max one paragraph). You do not repeat the words you used earlier, make it concise. Make it evident how your results answer the question you have asked.
- Compare your results with something in the literature (use PubMed, or literature indicated in UniProt, PDB or any datasets I learnt about). 1-2 articles enough. The point is that you
- show how your results relate to previous ones. Different scenarios: 'support', 'in contrast', 'in addition', 'gives insights into'.
- If possible, present a model. This can be a molecular 'mechanism', 'scenario', 'reason'. For example: 'A and B belong to the same family'; 'mutation XD likely changes the interaction in liver'; 'in disease T I have less phosphorylation sites or the level of phosphorylation decreased'; 'in this experiment, there is a problem with degradation of protein P.' etc.
- The model can be presented as: 'propose', 'suggest', 'indicate', 'hypothesize', or even 'support' or 'demonstrate'. If possible try to present a cartoon figure or scheme.

V. CONCLUSIONS

One paragraph, emphasizing the main finding and formulate a message for the future (perspective), how to continue from here, what kind of further studies are needed. You may have Conclusion as chapter 5.

VI. REFERENCES

The text should contain references, to previous results, data, methods, softwares, etc.. Format is unrestricted, but you should use these refs in the text, and list them at the end. I expect ca 15-20 refs listed.

Data availability statement: Preprocessed trajectories of the HP35 villin headpiece are kindly made available by Nagel et al. and downloadable at their GitHub repository <https://github.com/moldyn/HP35-DESRES>. All code necessary to replicate the analysis should in this work is available on GitHub at <https://github.com/miriamzara/MSM>.

-
- [1] Bowman, Gregory R. and Clementi, Cecilia and Huang, Xuhui, *Introduction to Markov State Modeling of Conformational Dynamics*, Journal of Chemical Theory and Computation, year 2026 volume 22. DOI: 10.1021/acs.jctc.6c00351.
 - [2] S. Piana, K. Lindorff-Larsen and D.E. Shaw, *Protein folding kinetics and thermodynamics from atomistic simulation*, Proc. Natl. Acad. Sci. U.S.A. 109 (44) 17845-17850 (2012), DOI: 10.1073/pnas.1201811109.
 - [3] J. Kubelka, T.K. Chiu, D.R. Davies et al., *Sub-microsecond protein folding*, J. Mol. Biol. 2006, 359, 546-553. DOI: 10.1016/j.jmb.2006.03.034.
 - [4] D. Nagel, S. Sartore, and G. Stock, *Selecting Features for Markov Modeling: A Case Study on HP35*, Journal of Chemical Theory and Computation 2023 19 (11), 3391-3405 DOI: 10.1021/acs.jctc.3c00240.
 - [5] D. Nagel, S. Sartore, and G. Stock, *Toward a Benchmark for Markov State Models: The Folding of HP35*, The Journal of Physical Chemistry Letters 2023 14 (31), 6956-6967 DOI: 10.1021/acs.jpcllett.3c01561 .
 - [6] A. Jain and G. Stock, *Hierarchical Folding Free Energy Landscape of HP35 Revealed by Most Probable Path Clustering* J. Phys. Chem. B, 118, 28, 7750-7760 (2014). DOI: 10.1021/jp410398a.
 - [7] G.R. Bowman, V.S. Pande and F. Noé, *An Introduction to Markov State Models and Their Application to Long Timescale Molecular Simulation* Advances in Experimental Medicine and Biology, Springer Dordrecht (2013) DOI: 10.1007/978-94-007-7606-7.
 - [8] M. Hoffmann, M.K. Scherer, T. Hempel et al., *Deeptime: a Python library for machine learning dynamical models from time series data*, Machine Learning: Science and

Technology, 2021, IOP Publishing. DOI: 10.1088/2632-2153/ac3de0.

- [9] A. Mardt, L. Pasquali, H. Wu and F. Noé, *VAMPnets for deep learning of molecular kinetics*, Nature Communications 9(1):1-11, 2018.
- [10] C. Ihmeyer and F. Noé, *Time-lagged autoencoders: deep learning of slow collective variables for molecular kinetics*. J. Chem. Phys., 148(24):241703, 2018. DOI: 10.1063/1.5011399.
- [11] F. Sittel, T. Filk and G. Stock, *Principal component analysis on a torus: Theory and application to protein dynamics* J Chem Phys. 147(24):244101, 2017. DOI: 10.1063/1.4998259.
- [12] G. Diez, D. Nagel and G. Stock, *Correlation-based feature selection to identify functional dynamics in proteins*. J. Chem. Theory Comput. 2022, 18, 5079-5088. DOI: 10.1021/acs.jctc.2c00337.
- [13] C.R. Schwantes and V.S. Pande, *Improvements in Markov State Model Construction Reveal Many Non-Native Interactions in the Folding of NTL9* J. Chem. Theory Comput. 2013, 9, 4, 2000-2009 DOI: 10.1021/ct300878a.
- [14] *PyMOL* The PyMOL Molecular Graphics System, Version 3.0 Schrödinger, LLC.

Appendix A: Supporting Information

S1. Derivation of PCA and tICA

The presentation of tICA in this section follows closely that of [13].

PCA and tICA are both linear dimensionality reduction techniques, i.e. they project the data onto a linear subspace of the original feature space. Let $\{\mathbf{x}(t)\}_{t=0}^{T-1}$ be a multidimensional, discrete time-series, where each snapshot is a column vector, of dimension D , corresponding to an arbitrary, vectorized representation of a protein conformation (e.g. a set of dihedral angles). Suppose the data has been sampled from a *stationary* and *time-reversible* stochastic process. Without loss of generality, I assume that the data have been centered, $\mathbb{E}[\mathbf{x}] = 0$. Using the bracket notation $\langle x | := \mathbf{x}^T$, $|x\rangle := \mathbf{x}$, I define the covariance matrix:

$$\mathbf{C}_0 := \mathbb{E}[|x\rangle \langle x|] \quad (\text{A1})$$

and the time-lagged covariance matrix:

$$\mathbf{C}_{0,\tau} := \mathbb{E}[|x(t)\rangle \langle x(t+\tau)|]. \quad (\text{A2})$$

For a stationary reversible process,

$$\mathbf{C}_{0,\tau} = \mathbf{C}_{\tau,0}^T = \mathbf{C}_{\tau,0},$$

so that the lagged covariance matrix may be denoted simply by $\mathbf{C}_\tau$.

If $|\hat{u}\rangle \langle \hat{u}|$ is a projection operator (i.e. $|\hat{u}\rangle$ is a unit vector), then:

$$\langle \hat{u} | \mathbf{C}_0 | \hat{u} \rangle = \mathbb{E}[\langle \hat{u} | x \rangle \langle x | \hat{u} \rangle] = \mathbb{E}[\langle \hat{u} | x \rangle^2] \quad (\text{A3})$$

is the variance of the projection of the data along $|\hat{u}\rangle$, and

$$\begin{aligned} \text{ACF}_{\hat{u}}(\tau) &= \frac{\langle \hat{u} | \mathbf{C}_\tau | \hat{u} \rangle}{\langle \hat{u} | \mathbf{C}_0 | \hat{u} \rangle} \\ &= \frac{\mathbb{E}[\langle \hat{u} | x(t) \rangle \langle \hat{u} | x(t+\tau) \rangle]}{\mathbb{E}[\langle \hat{u} | x \rangle^2]} \end{aligned} \quad (\text{A4})$$

its autocorrelation at lag time τ .

PCA seeks the direction of maximal variance. The first principal component is defined as

$$\begin{aligned} |\hat{u}_1\rangle &= \arg \max \langle \hat{u} | \mathbf{C}_0 | \hat{u} \rangle \\ &\text{subject to } \|\hat{u}\| = 1 \end{aligned} \quad (\text{A5})$$

this optimization problem can be solved introducing a Lagrange multiplier λ , and the stationarity condition yields

$$\mathbf{C}_0 |\hat{u}_i\rangle = \lambda_i |\hat{u}_i\rangle, \quad (\text{A6})$$

i.e. the projection vector is an eigenvector of the covariance matrix. The corresponding eigenvalue equals the variance of the data projected onto the principal component:

$$\langle \hat{u}_i | \mathbf{C}_0 | \hat{u}_i \rangle = \lambda_i \cdot \langle \hat{u}_i | \hat{u}_i \rangle = \lambda_i.$$

Then the second principal component is defined as the direction of maximal variance orthogonal to the first, giving the optimization problem:

$$\begin{aligned} |\hat{u}_2\rangle &= \arg \max \langle u | \mathbf{C}_0 | u \rangle \\ &\text{subject to } \langle u | u \rangle = 1 \\ &\text{and } \langle u | \hat{u}_1 \rangle = 0 \end{aligned} \quad (\text{A7})$$

which again can be solved with Lagrange multipliers. Iterating, this implies that the principal components form an orthonormal basis of eigenvectors of the covariance matrix, i.e. they diagonalize the covariance matrix. Dimensionality reduction ($d < D$) is achieved by taking the components along the first d principal components $\{|\hat{u}_1\rangle, \dots, |\hat{u}_d\rangle\}$.

$$z_i = \langle \hat{u}_i | x \rangle, \quad \mathbf{z} = U^T \mathbf{x}. \quad (\text{A8})$$

This is a linear projection onto an ortonormal basis:

$$|x\rangle \rightarrow |z\rangle := \sum_{i=1}^d \langle x|\hat{u}_i\rangle |\hat{u}_i\rangle$$

By construction, in the reduced space components are uncorrelated and ordered by decreasing variance:

$$\mathbb{E}[|z\rangle \langle z|] = \text{diag}(\lambda_1, \dots, \lambda_d) \quad (\text{A9})$$

tICA Unlike PCA, tICA [13] seeks directions $|u\rangle$ that maximize the autocorrelation Eq.(A4) at a fixed lag time τ . The lagtime τ is a free parameter. The autocorrelation is invariant under scaling of the projection vector $|u\rangle$. Differently from PCA, the normalization constraint is usually implemented as $\langle u|\mathbf{C}_0|u\rangle = 1$, as this choice simplifies the optimization problem. Formally, tICA solves the optimization problem:

$$\begin{aligned} |\hat{u}_1\rangle = & \text{argmax} \langle u|\mathbf{C}_\tau|u\rangle \\ \text{subject to } & \langle u|\mathbf{C}_0|u\rangle = 1 \end{aligned} \quad (\text{A10})$$

using the same procedure as for PCA, I can solve this with Lagrange multipliers, and show that

$$\mathbf{C}_\tau |\hat{u}_1\rangle = \lambda_1 \mathbf{C}_0 |\hat{u}_1\rangle, \quad (\text{A11})$$

The generalized eigenvalue λ_i is equal to the autocorrelation of the projected coordinate $\langle \hat{u}_i|x(t)\rangle$ at lag time τ . Assuming that this coordinate approximates a single relaxation mode of the dynamics, its autocorrelation decays exponentially, yielding

$$\lambda_i \approx e^{-\tau/t_i}, \quad (\text{A12})$$

where t_i is the corresponding relaxation timescale. In practice, the relation is approximate because a tIC may contain contributions from multiple dynamical modes.

Similarly as done with PCA, I can iterate this procedure and define the second tIC as the direction of maximal autocorrelation at lag time τ that is *uncorrelated* with the first:

$$\begin{aligned} |\hat{u}_2\rangle = & \text{argmax} \langle u|\mathbf{C}_\tau|u\rangle \\ \text{subject to } & \langle u|\mathbf{C}_0|u\rangle = 1 \\ \text{and } & \langle u|\mathbf{C}_0|\hat{u}_1\rangle = 0 \end{aligned} \quad (\text{A13})$$

the solution $|\hat{u}_1\rangle$ is again a solution of the same generalized eigenvalue problem as in Eq.(A10). generalized

eigenvalue problem. Dimensionality reduction ($d < D$) is achieved by taking the components of the data $\mathbf{x}(t)$ along the first d tICs $\{|\hat{u}_1\rangle, \dots, |\hat{u}_d\rangle\}$.

$$\mathbf{z}(t) = (\langle x(t)|\hat{u}_1\rangle, \dots, \langle x(t)|\hat{u}_d\rangle) \quad (\text{A14})$$

By construction, in the reduced space components z_i have unit variance, are uncorrelated at zero lag, and are ordered by decreasing autocorrelation at lag time τ .

$$\mathbb{E}[|z\rangle \langle z|] = \mathbb{I} \quad (\text{A15})$$

tICA as a diagonalization in the whitened coordinate space It is worth noting that unlike in PCA, the tIC vectors are generally not orthogonal under the standard Euclidean inner product. Instead, as I said, they satisfy the generalized orthogonality relation

$$\langle u_i|\mathbf{C}_0|u_j\rangle = \delta_{ij}.$$

i.e. they corresponding components z_i and z_j are *uncorrelated* at zero lag. Nevertheless, the tICA procedure can be reformulated as a standard eigenvalue problem in a transformed coordinate system. The key idea is to first apply a whitening transformation, which rescales and rotates the original data so that all directions have unit variance and are mutually uncorrelated. Define the ZCA whitening transformation

$$|y\rangle = \mathbf{C}_0^{-1/2} |x\rangle$$

where $\mathbf{C}_0^{-1/2}$ is symmetric and positive definite. By construction, the covariance matrix in the transformed space is the identity:

$$\mathbb{E}[|y\rangle \langle y|] = \mathbf{C}_0^{-1/2} \mathbf{C}_0 \mathbf{C}_0^{-1/2} = \mathbf{I}.$$

and time-lagged covariance matrix in the whitened space is

$$\hat{\mathbf{C}}_\tau = \mathbf{C}_0^{-1/2} \mathbf{C}_\tau \mathbf{C}_0^{-1/2}$$

$\hat{\mathbf{C}}_\tau$ is symmetric and can thus be diagonalized;

$$\hat{\mathbf{C}}_\tau |v_i\rangle = \lambda_i |v_i\rangle, \quad \langle v_i|v_j\rangle = \delta_{ij}.$$

If $|u_i\rangle$ solves the tICA problem (A10), then $|v_i\rangle = \mathbf{C}_0^{1/2} |u_i\rangle$ is an eigenvector of $\hat{\mathbf{C}}_\tau$ with the same eigenvalue.

Indeed, multiplying the generalized eigenvalue problem (A10) from the left by $\mathbf{C}_0^{-1/2}$ and defining $|v_i\rangle = \mathbf{C}_0^{1/2} |u_i\rangle$ yields

$$\hat{\mathbf{C}}_\tau |v_i\rangle = \lambda_i |v_i\rangle.$$

Consequently, the tICA components z_i correspond to the orthogonal projection of the whitened data $\mathbf{y}$ onto the orthonormal basis defined by $\{|v_i\rangle\}$:

$$z_i(t) = \langle x(t) | u_i \rangle = \langle x(t) | C_0^{-1/2} v_i \rangle = \langle y(t) | v_i \rangle.$$

Alternative generalized eigenvector normalizations for tICA

The construction presented above corresponds to the original formulation of tICA, in which the covariance matrix of the transformed data is the identity (A15). HoIver, alternative rescaling of the transformed coordinates have been used in the literature and are usually implemented in softwares to endow Euclidean distances in the reduced space with a more direct kinetic interpretation. These can be particularly useful when the tICA coordinates are subsequently clustered to construct a Markov state model, as I do in this work.

The most commonly used choice is the *kinetic map* in which each tIC component z_i is multiplied by its corresponding eigenvalue,

$$\tilde{z}_i(t) = \lambda_i z_i(t). \quad (\text{A16})$$

With this normalization, the coordinates representing the slo1st processes "light more", and a unit step in those directions in the reduced space is assigned a larger distance. An alternative is the *commute map*, where each coordinate is rescaled according to its implied timescale,

$$\tilde{z}_i(t) = \sqrt{\frac{t_i}{2}} z_i(t), \quad t_i = -\frac{\tau}{\ln \lambda_i}. \quad (\text{A17})$$

With this normalization, Euclidean distances approximate commute distances, i.e. the expected time required to travel from one configuration to another and back. This is a variation on the same idea behind the kinetic map, being eigenvalues and timescales connected by (A12).

S2. Spectral decomposition of the transition matrix of a Markov chain **TO COMPLETE**

The dynamics encoded by a Markov state model can be understood through the eigendecomposition of its transition matrix $\mathbf{T}$. For consistency with the notation used in the Python package deeptime, let us consider the transition matrix $\mathbf{T}$ as a row-stochastic matrix, i.e. T_{ij} is the probability of transitioning from state i to state j in one lag step:

$$\sum_j T_{ij} = 1, \quad T_{ij} \geq 0.$$

Probability distributions are represented as row vectors, and evolve as:

$$\mathbf{p}(n) = \mathbf{p}(0)\mathbf{T}^n,$$

where $\mathbf{p}(0)$ is the initial distribution and $\mathbf{p}(n)$ is the distribution after n lag steps. The transition matrix admits a spectral decomposition of the kind:

$$\mathbf{T} = \sum_i \lambda_i \mathbf{r}_i \mathbf{l}_i^T,$$

where λ_i are the eigenvalues, and $\mathbf{r}_i$ and $\mathbf{l}_i$ are the corresponding right and left eigenvectors, satisfying:

$$\mathbf{T} \mathbf{r}_i = \lambda_i \mathbf{r}_i, \quad \mathbf{l}_i^T \mathbf{T} = \lambda_i \mathbf{l}_i^T.$$

The probability distribution after n lag times can therefore be written as

$$\mathbf{p}(n) = \mathbf{p}(0)\mathbf{T}^n = \sum_i c_i \cdot \lambda_i^n \cdot \mathbf{l}_i,$$

where the coefficients $c_i = \mathbf{p}(0)\mathbf{r}_i$ depend on the initial condition. The dominant eigenvalue is $\lambda_1 = 1$ and corresponds to the equilibrium distribution. All other modes decay exponentially with characteristic relaxation times

$$t_i = \frac{\tau}{\ln \lambda_i},$$

where τ is the MSM lag time.

The left eigenvectors therefore describe the spatial structure of the relaxation processes occurring on the corresponding timescales. States with eigenvector components of opposite sign exchange probability during the relaxation process, while the magnitude of each component quantifies the degree to which that state participates in the process. The instantaneous probability flux associated with a mode can be visualized from the change in the probability distribution after one lag step,

$$\Delta \mathbf{p} = \mathbf{p} \mathbf{T} - \mathbf{p} = \text{TODO : complete}$$

For reversible Markov models obeying detailed balance, left and right eigenvectors contain equivalent information and are related through the equilibrium population matrix $\mathbf{\Pi}$,

$$\mathbf{l}_i \mathbf{\Pi}, \mathbf{r}_i.$$

Consequently, the right eigenvectors may be interpreted as equilibrium-population-weighted versions of the relaxation modes.

Appendix B: Supplementary Figures




FIG. SF1: Flowchart of the MSM construction pipeline used in this work.




FIG. SF2: **tICA projections obtained from the dihedral-angle feature set using kinetic-map scaling.** Columns correspond to different tICA lag times spanning four orders of magnitude ($10, 10^2, 10^3, 10^5$ frames; values in μ s are reported above each panel). The top row shows the free-energy landscape projected onto the first two tICs, while the loIr panels report the first seven tIC time series (left) and their one-dimensional free-energy profiles (right) for the first 40μ s of simulation. A useful tICA projection should separate slowly interconverting conformations into distinct regions of configuration space and produce components exhibiting metastable, jump-like dynamics. Very short lag times lead to projections dominated by fast fluctuations, whereas excessively large lag times wash out the dynamical structure and yield nearly unimodal projections.

Intermediate lag times ($\tau \sim 1\text{--}2 \mu$ s) provide the clearest separation of metastable states and the most pronounced residence/jump behaviour in the leading components. Based on this analysis, a lag time of approximately 2μ s and the first five tICs were retained for subsequent MSM construction.




FIG. SF3: tICA projections obtained from the contact-distance feature set using kinetic-map scaling.

Compared to the dihedral-angle representation, the contact-distance features yield tICA coordinates with less clearly separated free-energy minima and less pronounced metastable, jump-like dynamics. This suggests a laker separation of slow conformational processes in the resulting low-dimensional projection. A lag time of $\tau = 10^3 \text{ frames} (\approx 0.2 \mu\text{s})$ was selected to maintain consistency with the analysis performed on the dihedral feature set. At this lag time, the free-energy profiles become largely unimodal beyond the first few tICs, indicating that most of the slow dynamical information is concentrated in the leading components. Although the first three tICs appear to capture the dominant slow processes, as a more conservative choice the first five components were retained for subsequent MSM construction.




FIG. SF4: **Comparison of PCA projections obtained from the contact-distance feature set and from the dihedrals feature set.** In contrast to tICA, the PCA projections provide a less clear separation of metastable conformational states. Free-energy profiles exhibit Iaker clustering, the one-dimensional free-energy profiles are predominantly unimodal, and the principal-component time series do not display a pronounced residence/jump behavior, with exception of the first PC. Consistent with the tICA results, the dihedral-angle representation yields a more structured low-dimensional projection than the contact-distance features. For subsequent MSM construction I retain the first 7 components for the dihedrals set and the first 5 for the contact-distance set.




FIG. SF5: **Free-energy landscapes and structural observables in reduced coordinate spaces.** **Top row:** free-energy surfaces ΔF projected onto the first two PCA or tICA components, calculated as $F(x, y) = -\ln p(x, y)$, $\Delta F(x, y) = F(x, y) - F_{min}$. **Middle row:** projections coloured by the fraction of native contacts Q . **Bottom row:** projections coloured by the RMSD to the crystal structure (normalized). Each column corresponds to a different combination of feature set (distances or dihedral angles) and dimensionality reduction method (PCA or tICA). The colouring reveals the extent to which the reduced coordinates correlate with physically meaningful progress variables of folding, namely native-contact formation and structural similarity to the crystal structure.